

Swiss Institute of  
Bioinformatics

# Advanced R

## September 2024

Frédéric Schütz      (`Frederic.Schutz@sib.swiss`)  
Frédéric Burdet      (`Frederic.Burdet@sib.swiss`)



# **An introduction (or reminder) about R data structures**

All objects in R have a **type**, which describes the type of data stored in the object.

## *The typeof ( ) command*

To find the type of any object:

```
typeof (object)
```

```
> typeof( c(1,2,3) )  
[1] "double"  
> typeof( c("a", "b", "c") )  
[1] "character"
```

## *Possible types in R*

- logical
- integer
- double
- closure
- builtin
- special
- complex
- character
- raw
- list

(and a few others)

Sometimes, we also talk about the **mode**, a simplified version of types.

## *The mode ( ) command*

To find the mode of any object:

```
mode(object)
```

```
> typeof( c(1,2,3) )  
[1] "double"  
> typeof( c("a", "b", "c") )  
[1] "character"  
  
> mode( c(1,2,3) )  
[1] "numeric"  
> mode( c("a", "b", "c") )  
[1] "character"
```

## *Possible types in R*

### **Type**

- logical
- integer
- double
- closure
- builtin
- special
- complex
- character
- raw
- list



## *Possible types and modes in R*

### **Type**

- logical
- integer
- double
- closure
- builtin
- special
- complex
- character
- raw
- list

### **Mode**

logical  
numeric  
numeric  
function  
function  
function  
complex  
character  
raw  
list

## *An explanation for some abstruse error messages in R*

```
> x <- 1
> print(c(class(x), mode(x), typeof(x)))
[1] "numeric" "numeric" "double"
> x <- letters
> print(c(class(x), mode(x), typeof(x)))
[1] "character" "character" "character"
> x <- TRUE
> print(c(class(x), mode(x), typeof(x)))
[1] "logical" "logical" "logical"
> print(c(class(x), mode(x), typeof(x)))
[1] "logical" "logical" "logical"
> x <- cars[1]
> print(c(class(x), mode(x), typeof(x)))
[1] "data.frame" "list"      "list"
> x <- cars[[1]]
> print(c(class(x), mode(x), typeof(x)))
[1] "numeric" "numeric" "double"
> x <- matrix(cars)
> print(c(class(x), mode(x), typeof(x)))
[1] "matrix" "array"   "list"    "list"
> x <- new.env()
> print(c(class(x), mode(x), typeof(x)))
[1] "environment" "environment" "environment"
> x <- ls
> print(c(class(x), mode(x), typeof(x)))
[1] "function" "function" "closure"
```

## *An explanation for some abstruse error messages in R*

- |           |          |
|-----------|----------|
| • logical | logical  |
| • integer | numeric  |
| • double  | numeric  |
| • closure | function |

```
> f <- function() {}
```

```
> f$a
```

```
Error in f$a : object of type 'closure' is not subsettable
```

# Vectors

The simplest way to store data into R is the vector, which contains an ordered collection of objects **of the same type**:

```
> x <- c(1, 2, 3, 4); x  
[1] 1 2 3 4  
> typeof(x); mode(x)  
[1] "double"  
[1] "numeric"
```

The simplest way to store data into R is the vector, which contains an ordered collection of objects **of the same type**:

```
> x <- c(1, 2, 3, 4); x
```

```
[1] 1 2 3 4
```

```
> typeof(x); mode(x)
```

```
[1] "double"
```

```
[1] "numeric"
```

```
> x <- c(TRUE, FALSE, TRUE, TRUE)
```

```
> typeof(x); mode(x)
```

```
[1] "logical"
```

```
[1] "logical"
```

*What happens if I store several types of objects in a vector ?*

```
> x <- c(1, 2, TRUE, 3); x
```

*What happens if I store several types of objects in a vector ?*

```
> x <- c(1, 2, TRUE, 3); x  
[1] 1 2 1 3  
> typeof(x)  
[1] "double"
```

R will convert the objects to the type that is able to accommodate all of them.



*What happens if I store several types of objects in a vector ?*

```
> x <- c(1, 2, TRUE, 3); x
```

```
[1] 1 2 1 3
```

```
> typeof(x)
```

```
[1] "double"
```

```
> x <- c(1, 2, "true", 4); x
```

*What happens if I store several types of objects in a vector ?*

```
> x <- c(1, 2, TRUE, 3); x
```

```
[1] 1 2 1 3
```

```
> typeof(x)
```

```
[1] "double"
```

```
> x <- c(1, 2, "true", 4); x
```

```
[1] "1"      "2"      "true"   "4"
```

```
> typeof(x)
```

```
[1] "character"
```

*With some surprises...*

```
> x <- c("a", TRUE, 3); x  
[1] "a"      "TRUE"    "3"  
> typeof(x)  
[1] "character"
```

*With some surprises...*

```
> x <- c("a", TRUE, 3); x
```

```
[1] "a"      "TRUE"   "3"
```

```
> typeof(x)
```

```
[1] "character"
```

```
> x <- c("a", c(TRUE, 3)); x
```

```
[1] "a" "1" "3"
```

```
> typeof(x)
```

```
[1] "character"
```

Logical values (TRUE/FALSE) can easily be converted to numeric value (0/1) and back, as in most programming languages:

```
> as.numeric( c(FALSE, TRUE) )
```

```
[1] 0 1
```

```
> as.logical( c(0,1) )
```

```
[1] FALSE TRUE
```

```
> c(FALSE, 0, TRUE)
```

```
[1] 0 0 1
```

## *Counting values using logicals*

This is very useful, for example for counting purposes.

Example: count **the number of elements** of the vector `data` that are larger than zero:

```
> data <- rnorm(10)
> data
[1] -0.61518461 -0.62574053  1.21586046 -1.42627945
[5]  0.06749257  0.59811401  0.25876230 -0.45936110
[9] -1.83171441  0.28693148
> data > 0
[1] FALSE FALSE  TRUE FALSE  TRUE
[6]  TRUE  TRUE FALSE FALSE  TRUE
> sum(data > 0)
[1] 5
```

Example: count **the proportion of elements** of the vector `data` that are larger than zero:

```
> data <- rnorm(10)
> data
[1] -0.61518461 -0.62574053  1.21586046 -1.42627945
[5]  0.06749257  0.59811401  0.25876230 -0.45936110
[9] -1.83171441  0.28693148
> data > 0
[1] FALSE FALSE  TRUE FALSE  TRUE
[6]  TRUE  TRUE FALSE FALSE  TRUE
> mean(data > 0)
[1] 0.5
```

*Why do the two selection commands return different results ?*

```
> vector <- 1:10
```

```
> vector[ c(0,1) ]
```

```
[1] 1
```

```
> vector[ c(FALSE, TRUE) ]
```

```
[1] 2 4 6 8 10
```



## *Difference between logical and numeric*

```
> vector <- 1:10  
  
> vector[ c(0,1) ]  
[1] 1
```

This selects elements 0 (which does not exist) and 1 (=1)

```
> vector[ c(FALSE,TRUE) ]  
[1] 2 4 6 8 10
```

This applies to each element in turn; since the logical vector is not long enough, it is recycled to cover the full vector. At the end, only elements at even positions are selected.

## *Difference between logical and numeric*

```
> vector <- 1:10  
  
> vector[ c(0,1) ]  
[1] 1  
  
> vector[ c(FALSE, TRUE) ]  
[1] 2 4 6 8 10
```

In contrast to other programming languages, logical and numeric types can not be freely exchanged !

## *How to use this feature to your advantage*

This behaviour can lead to bugs in your code. But if done correctly, it can also help you.

## *How to use this feature to your advantage*

This behaviour can lead to bugs in your code. But if done correctly, it can also help you.

The following command is a simple way to select all elements at even positions in the vector.

(for example, sampling one data point out of two)

```
> vector[ c(FALSE,TRUE) ]  
[1]  2  4  6  8 10
```

*What could possibly go wrong ?*

```
> sample(1:10, 10, replace=T)
[1]  8  4  9  1 10  3  9  6  3  9
```

*When T does not mean TRUE*

```
> A <- "a"; B <- "b"; C <- "c"; T <- "t"
```

```
> sample(1:10, 10, replace=T)
```

```
> A <- "a"; B <- "b"; C <- "c"; T <- "t"
```

```
> sample(1:10, 10, replace=T)
```

```
Error in sample(1:10, 10, replace = T) : invalid 'replace' argument
```

```
> A <- "a"; B <- "b"; C <- "c"; T <- "t"
```

```
> sample(1:10, 10, replace=T)
```

```
Error in sample(1:10, 10, replace = T) : invalid 'replace' argument
```

‘T’ and ‘F’ can be freely redefined by the user, something impossible with the full form:

```
> TRUE <- "t"
```

```
Error in TRUE <- "t" : invalid (do_set) left-hand side to assignment
```

This will yield an error, or even worse...



*If you are really vicious...*

```
> T <- FALSE
```

```
> sample(1:10, 10, replace=T)
```

```
[1]  7  6  3  4 10  1  8  5  9  2
```

*If you are really vicious...*

```
> T <- FALSE  
  
> sample(1:10, 10, replace=T)  
[1] 7 6 3 4 10 1 8 5 9 2
```

Or, more likely:

```
> T <- complicated_function(many, many, complicated, arguments, and  
                             the, function, returns, FALSE, in, the,  
                             end )  
  
> sample(1:10, 10, replace=T)  
[1] 7 6 3 4 10 1 8 5 9 2
```

**Attributes** are arbitrary labels attached to the R objects.

**Attributes** are arbitrary labels attached to the R objects.

```
> x <- rnorm(10)
> attributes(x)
NULL
```

**Attributes** are arbitrary labels attached to the R objects.

```
> x <- rnorm(10)
> attributes(x)
NULL
> attr(x, "mylabel") <- "Random normal data"
> attr(x, "mylabel")
[1] "Random normal data"
```

**Attributes** are arbitrary labels attached to the R objects.

```
> x <- rnorm(10)
> attributes(x)
NULL
> attr(x, "mylabel") <- "Random normal data"
> attr(x, "mylabel")
[1] "Random normal data"
> attributes(x)
$mylabel
[1] "Random normal data"
```

## *Some important attributes in R*

- **names:** allows naming the components of an object
- **class:** a label attached to the object, which indicates how actions can be performed on the object
- **dim:** the dimensions of the objects (e.g. for a matrix or an array)

## *Adding a class*

```
> x <- rnorm(10)
> attributes(x)
NULL
> attr(x, "mylabel") <- "Random normal data"
> attr(x, "mylabel")
[1] "Random normal data"
> attributes(x)
$mylabel
[1] "Random normal data"

> class(x) <- "randomdata"
```



## *Adding a class*

```
> x <- rnorm(10)
> attributes(x)
NULL
> attr(x, "mylabel") <- "Random normal data"
> attr(x, "mylabel")
[1] "Random normal data"
> attributes(x)
$mylabel
[1] "Random normal data"

> class(x) <- "randomdata"
> attr(x, "class") <- "randomdata"    # equivalent
```

## *Adding a class*

```
> x <- rnorm(10)
> attributes(x)
NULL
> attr(x, "mylabel") <- "Random normal data"
> attr(x, "mylabel")
[1] "Random normal data"
> attributes(x)
$mylabel
[1] "Random normal data"

> class(x) <- "randomdata"
> attr(x, "class") <- "randomdata"    # equivalent
> class(x)
[1] "randomdata"
```

## *Adding names*

```
> names(x) <- LETTERS[1:10]
```

```
> x
```

A	B	C	D	E
-0.93205027	-0.16194958	0.26727310	-0.07427123	1.54048877
F	G	H	I	J
-0.63579513	0.27141749	-2.03039854	-2.52658864	1.02263626

```
attr(,"mylabel")
```

```
[1] "Random normal data"
```

```
attr(,"class")
```

```
[1] "randomdata"
```

```
> attributes(x)
```

```
$mylabel
```

```
[1] "Random normal data"
```

```
$class
```

```
[1] "randomdata"
```

```
$names
```

```
[1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J"
```

## *Adding names*

```
> names(x) <- LETTERS[1:10]
```

```
> x
```

A	B	C	D	E
-0.93205027	-0.16194958	0.26727310	-0.07427123	1.54048877
F	G	H	I	J
-0.63579513	0.27141749	-2.03039854	-2.52658864	1.02263626

```
attr(,"mylabel")
```

```
[1] "Random normal data"
```

```
attr(,"class")
```

```
[1] "randomdata"
```

```
> attributes(x)
```

```
$mylabel
```

```
[1] "Random normal data"
```

```
$class
```

```
[1] "randomdata"
```

```
$names
```

```
[1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J"
```

## *Why names are important*

Names allow you to create **lookup tables**

```
> names(x) <- LETTERS[1:10]
```

```
> x
```

A	B	C	D	E
-0.93205027	-0.16194958	0.26727310	-0.07427123	1.54048877
F	G	H	I	J
-0.63579513	0.27141749	-2.03039854	-2.52658864	1.02263626

```
> x["B"]
```

```
[1] -0.16194958
```

## Names allow you to create **lookup tables**

```
> names(x) <- LETTERS[1:10]
```

```
> x
```

A	B	C	D	E
-0.93205027	-0.16194958	0.26727310	-0.07427123	1.54048877
F	G	H	I	J
-0.63579513	0.27141749	-2.03039854	-2.52658864	1.02263626

```
> x["B"]
```

```
[1] -0.16194958
```

```
> x[2]
```

```
[1] -0.16194958
```

# Matrices and arrays

Matrices (in 2D) and arrays (in 2D or more) are extensions of vectors, where two or more dimensions are specified.

```
> m <- matrix(1:30, ncol=6)
> m
```

	[, 1]	[, 2]	[, 3]	[, 4]	[, 5]	[, 6]
[1, ]	1	6	11	16	21	26
[2, ]	2	7	12	17	22	27
[3, ]	3	8	13	18	23	28
[4, ]	4	9	14	19	24	29
[5, ]	5	10	15	20	25	30

```
> m[1, 3]
[1] 11
```



In fact, a matrix (or array) is stored as a vector (column by column) with additional information about its dimensions.

```
> m <- matrix(1:30, ncol=6)
```

```
> m
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	1	6	11	16	21	26
[2,]	2	7	12	17	22	27
[3,]	3	8	13	18	23	28
[4,]	4	9	14	19	24	29
[5,]	5	10	15	20	25	30

```
> as.vector(m)
```

[1]	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
[16]	17	18	19	20	21	22	23	24	25	26	26	27	28	29	30

Internally, matrices are just vectors, with indications of dimensions.

```
> m <- matrix(1:30, ncol=6)
```

```
> m
```

	[, 1]	[, 2]	[, 3]	[, 4]	[, 5]	[, 6]
[1,]	1	6	11	16	21	26
[2,]	2	7	12	17	22	27
[3,]	3	8	13	18	23	28
[4,]	4	9	14	19	24	29
[5,]	5	10	15	20	25	30

```
> as.vector(m)
```

[1]	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
[16]	17	18	19	20	21	22	23	24	25	26	26	27	28	29	30

So you can also access matrices as if they were vectors.

```
> m <- matrix(1:30, ncol=6)

> m[11]; m[1,3]           # Equivalent
[1] 11
[1] 11
```

They have both a **length** and **dimensions**.

```
> m <- matrix(1:30, ncol=6)

> m[11]; m[1,3]           # Equivalent
[1] 11
[1] 11

> dim(m)
[1] 5 6

> length(m)
[1] 30
```

Arrays are constructed in a similar way.

```
> a <- 1:24  
> array(a, dim=c(4,3,2))  
, , 1
```

	[,1]	[,2]	[,3]
[1,]	1	5	9
[2,]	2	6	10
[3,]	3	7	11
[4,]	4	8	12

```
, , 2
```

	[,1]	[,2]	[,3]
[1,]	13	17	21
[2,]	14	18	22
[3,]	15	19	23
[4,]	16	20	24

## *Transforming a vector into a matrix*

You must specify the dimensions and change the class.

```
> a <- 1:30
> attr(a, "dim") <- c(5,6)
> class(a) <- "matrix"
> a
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	1	6	11	16	21	26
[2,]	2	7	12	17	22	27
[3,]	3	8	13	18	23	28
[4,]	4	9	14	19	24	29
[5,]	5	10	15	20	25	30

## *Creating a matrix row by row*

A matrix can also be created row by row, using the `byrow` parameter.

However, it will still be stored column by column.

```
> m <- matrix(1:30, ncol=6, byrow=TRUE); m
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	1	2	3	4	5	6
[2,]	7	8	9	10	11	12
[3,]	13	14	15	16	17	18
[4,]	19	20	21	22	23	24
[5,]	25	26	27	28	29	30

```
> as.vector(m)
```

[1]	1	7	13	19	25	2	8	14	20	26	3	9	15	21	27
[16]	4	10	16	22	28	5	11	17	23	29	6	12	18	24	30

## *Type of elements in a matrix*

Since they are vectors, all elements of matrices must be of the same type:

```
> m <- matrix(1:30, ncol=6)
> typeof(m)
[1] "integer"
> m[3,3] <- "a"
> m
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	"1"	"6"	"11"	"16"	"21"	"26"
[2,]	"2"	"7"	"12"	"17"	"22"	"27"
[3,]	"3"	"8"	"a"	"18"	"23"	"28"
[4,]	"4"	"9"	"14"	"19"	"24"	"29"
[5,]	"5"	"10"	"15"	"20"	"25"	"30"

```
> typeof(m)
[1] "character"
```



```
selectcolumns <- function( m, cols, rows ) {  
  m1 <- m [, cols]  
  m2 <- m1[rows, ]  
  m2  
}  
  
nrows <- 20  
m1 <- data.frame( a=runif(nrows), b=runif(nrows), c=runif(nrows) )  
row.names(m1) <- paste0("row", 1:nrow(m1))  
  
cols <- c("b", "c")  
rows <- c("row10", "row12")  
  
> selectcolumns(m1, cols, rows)  
           b           c  
row10 0.8578518 0.2864960  
row12 0.3767570 0.7874534
```

```
selectcolumns <- function( m, cols, rows ) {  
  m1 <- m [, cols]  
  m2 <- m1[rows, ]  
  m2  
}  
  
nrows <- 20  
m1 <- data.frame( a=runif(nrows), b=runif(nrows), c=runif(nrows) )  
row.names(m1) <- paste0("row", 1:nrow(m1))  
  
cols <- "b"  
rows <- c("row10", "row12")  
  
> selectcolumns(m1, cols, rows)
```

```
selectcolumns <- function( m, cols, rows ) {  
  m1 <- m [, cols]  
  m2 <- m1[rows, ]  
  m2  
}  
  
nrows <- 20  
m1 <- data.frame( a=runif(nrows), b=runif(nrows), c=runif(nrows) )  
row.names(m1) <- paste0("row", 1:nrow(m1))  
  
cols <- "b"  
rows <- c("row10", "row12")  
  
> selectcolumns(m1, cols, rows)  
Error in m1[rows, ] : incorrect number of dimensions
```

## *Subsetting a matrix*

```
> m <- matrix(1:6, nrow=2)
```

```
> m
```

	[,1]	[,2]	[,3]
[1,]	1	3	5
[2,]	2	4	6

## *Subsetting a matrix*

```
> m <- matrix(1:6, nrow=2)
```

```
> m
```

	[,1]	[,2]	[,3]
[1,]	1	3	5
[2,]	2	4	6

```
> m[,1:2]
```

	[,1]	[,2]
[1,]	1	3
[2,]	2	4

... yields a matrix.

## *Subsetting a matrix*

```
> m <- matrix(1:6, nrow=2)
```

```
> m
```

	[,1]	[,2]	[,3]
[1,]	1	3	5
[2,]	2	4	6

```
> m[,1]
```

## *Subsetting a matrix*

```
> m <- matrix(1:6, nrow=2)
```

```
> m
```

```
      [,1] [,2] [,3]  
[1,]    1    3    5  
[2,]    2    4    6
```

```
> m[,1]
```

```
[1] 1 2
```

... yields a **vector** (instead of a 2 x 1 matrix).

By default, R removes all dimensions  
that it deems not useful !

To avoid this, use the `drop=FALSE` option to the matrix subsetting:

```
> m[,1]
[1] 1 2
```

```
> m[, 1, drop=FALSE]
      [,1]
[1,]     1
[2,]     2
```

```
> m[1, , drop=FALSE]
      [,1] [,2] [,3]
[1,]     1     3     5
```



It is not possible to set `drop=FALSE` as the default mode.

Doing this would mean that accessing one element in a matrix would return a 1x1 matrix:

```
> m[2,3, drop=FALSE]  
      [,1]  
[1,]      6
```

... which is almost certainly not what you want.

## *Another possible consequence*

```
> head(data1, 3)
  identifier    var1    var2
1         3862 0.87207 -2.0105
2         1577 0.01075  0.1970
3         5150 1.28249 -0.4650
> head(data2, 3)
  identifier    var3    var4
1         3862 0.1383 -2.0165
2         1577 2.3219  0.6855
3         5150 0.6865  0.7783
> data <- cbind( data1[, c("var1", "var2")],
                  data2[, c("var3", "var4")],
                  data1[, "identifier"] )
```

## *Matrices converted to vectors lose their names !*

```
> head(data1, 3)
  identifier    var1    var2
1         3862 0.87207 -2.0105
2         1577 0.01075  0.1970
3         5150 1.28249 -0.4650

> head(data2, 3)
  identifier    var3    var4
1         3862 0.1383 -2.0165
2         1577 2.3219  0.6855
3         5150 0.6865  0.7783

> data <- cbind( data1[, c("var1", "var2")],
                 data2[, c("var3", "var4")],
                 data1[, "identifier"] )

> head(data, 3)
      var1    var2    var3    var4 data1[, "identifier"]
1 0.87207 -2.01057 0.13836 -2.0165                3862
2 0.01075  0.19709 2.32192  0.6855                1577
3 1.28249 -0.46507 0.68659  0.7783                5150
```

## *Matrices converted to vectors lose their names !*

```
> head(data1, 3)
  identifier    var1    var2
1         3862 0.87207 -2.0105
2         1577 0.01075  0.1970
3         5150 1.28249 -0.4650

> head(data2, 3)
  identifier    var3    var4
1         3862 0.1383 -2.0165
2         1577 2.3219  0.6855
3         5150 0.6865  0.7783

> data <- cbind( data1[, c("var1", "var2")],
                 data2[, c("var3", "var4")],
                 data1[, "identifier", drop=FALSE] )

> head(data, 3)
      var1    var2    var3    var4 identifier
1  0.87207 -2.01057 0.13836 -2.0165        3862
2  0.01075  0.19709 2.32192  0.6855        1577
3  1.28249 -0.46507 0.68659  0.7783        5150
```

*What happens if I store several types of objects in a vector ?*

```
> x <- c(1, "a", c); x
```

*What happens if I store several types of objects in a vector ?*

```
> x <- c(1, "a", c); x
```

*What happens if I store several types of objects in a vector ?*

```
> x <- c(1, "a", c); x
```

```
[[1]]
```

```
[1] 1
```

```
[[2]]
```

```
[1] "a"
```

```
[[3]]
```

```
function (...) .Primitive("c")
```

```
> typeof(x)
```

```
[1] "list"
```

Lists allow the storage of several objects (with different types) in a single R object.

```
> mylist <- list(ages=c(21, 32, 41, 45),  
                height=c(180, 176, 156, 165),  
                sex=c("M", "M", "F", "M") )
```



Lists allow the storage of several objects (with different types) in a single R object.

```
> mylist <- list(ages=c(21, 32, 41, 45),  
                 height=c(180, 176, 156, 165),  
                 sex=c("M", "M", "F", "M") )  
  
> mylist  
$ages  
[1] 21 32 41 45  
  
$height  
[1] 180 176 156 165  
  
$sex  
[1] "M" "M" "F" "M"
```

Lists allow the storage of several objects (with different types) in a single R object.

```
> mylist <- list(ages=c(21, 32, 41, 45),  
                 height=c(180, 176, 156, 165),  
                 sex=c("M", "M", "F", "M") )  
  
> mylist  
$ages  
[1] 21 32 41 45  
  
$height  
[1] 180 176 156 165  
  
$sex  
[1] "M" "M" "F" "M"  
  
> class(mylist); typeof(mylist)  
[1] "list"  
[1] "list"
```

**When accessing a list element,  
what is the difference between**

**`mylist[1]`**

**and**

**`mylist[[1]]` ?**

The objects can be accessed either using their rank, or by their name.

[  $\times$  ] returns part (one element) of the list

[ [  $\times$  ] ] returns what is inside this element

The objects can be accessed either using their rank, or by their name.

`[x]` returns part (one element) of the list

`[ [x] ]` returns what is inside this element

```
> mylist[1]
```

The objects can be accessed either using their rank, or by their name.

`[x]` returns part (one element) of the list

`[ [x] ]` returns what is inside this element

```
> mylist[1]
$ages
[1] 21 32 41 45
> typeof(mylist[1])
[1] "list"
```

The objects can be accessed either using their rank, or by their name.

`[x]` returns part (one element) of the list

`[ [x] ]` returns what is inside this element

```
> mylist[1]
$ages
[1] 21 32 41 45
> typeof(mylist[1])
[1] "list"

> mylist[[1]]
[1] 21 32 41 45
> typeof(mylist[[1]])
[1] "double"
```

The objects can be accessed either using their rank, or by their name.

`[x]` returns part (one element) of the list

`[ [x] ]` returns what is inside this element

```
> mylist[1]
$ages
[1] 21 32 41 45
> typeof(mylist[1])
[1] "list"

> mylist[[1]]
[1] 21 32 41 45
> typeof(mylist[[1]])
[1] "double"

> mylist$height
[1] 180 176 156 165
```



- **Atomic vectors:**  
an ordered collection of data  
of the **same type**
- **Lists:**  
an ordered collection of data  
that can be of **different types**.

# Data frames

Data frames are usually the preferred method for working with datasets that consists of several observations (rows) on several variables (columns).

4 observations

```
> data <- as.data.frame( mylist )  
> data  
  ages height sex  
1   21   180  M  
2   32   176  M  
3   41   156  F  
4   45   165  M
```

3 variables

## *Data frames*

```
> data <- as.data.frame( mylist )  
> data  
  ages height sex  
1   21   180  M  
2   32   176  M  
3   41   156  F  
4   45   165  M
```

## *Data frames*

```
> data <- as.data.frame( mylist )  
> data  
  ages height sex  
1   21   180  M  
2   32   176  M  
3   41   156  F  
4   45   165  M  
> class(data) ; typeof(data)
```

*Data frames are lists that look like matrices*

Data frames are actually lists in R.

```
> data <- as.data.frame( mylist )
> data
  ages height sex
1   21   180  M
2   32   176  M
3   41   156  F
4   45   165  M
> class(data) ; typeof(data)
[1] "data.frame"
[1] "list"
```

*Data frames are lists that look like matrices*

They are easier to use than lists: you can access the elements as in a matrix all elements, since they all have the same length

```
> data <- as.data.frame( mylist )  
> data  
  ages height sex  
1   21   180  M  
2   32   176  M  
3   41   156  F  
4   45   165  M  
> data[2,2]  
[1] 176
```

## *Data frames are lists that look like matrices*

They are easier to use than lists: you can access the elements as in a matrix all elements, since they all have the same length

They are more flexible than matrices, as they allow columns of different types, while still making them easy to access.

```
> data <- as.data.frame( mylist )  
> data  
  ages height sex  
1   21   180  M  
2   32   176  M  
3   41   156  F  
4   45   165  M  
> typeof(data[,1]); typeof(data[,3])  
[1] "double"  
[1] "character"
```



## *Data frames vs lists*

To convert a list into a matrix, one only needs to:

- change the class to `data.frame`
- give (unique) names to the rows by setting the `row.names` attribute

```
> class(mylist) <- "data.frame"
```

## *Data frames vs lists*

To convert a list into a matrix, one only needs to:

- change the class to `data.frame`
- give (unique) names to the rows by setting the `row.names` attribute

```
> class(mylist) <- "data.frame"  
> mylist  
[1] ages    height sex  
<0 rows> (or 0-length row.names)
```

## *Data frames vs lists*

To convert a list into a matrix, one only needs to:

- change the class to `data.frame`
- give (unique) names to the rows by setting the `row.names` attribute

```
> class(mylist) <- "data.frame"
> mylist
[1] ages    height sex
<0 rows> (or 0-length row.names)
> row.names(mylist) <- 1:length(mylist[[1]])
> mylist
  ages height sex
1   21   180   M
2   32   176   M
3   41   156   F
4   45   165   M
```

## *Data frames: accessing columns*

```
> data
  ages height sex
1   21   180  M
2   32   176  M
3   41   156  F
4   45   165  M
```

```
> data[1]
  ages
1   21
2   32
3   41
4   45
```

Columns can be accessed just like a list  
The result is a **single-column dataframe**.

## *Data frames: accessing columns*

```
> data
  ages height sex
1   21   180   M
2   32   176   M
3   41   156   F
4   45   165   M

> data[,1]
[1] 21 32 41 45

> data[[1]]
[1] 21 32 41 45
```

Alternatively, it is possible to access the **content** of a given column, yielding a vector.

Columns can also be accessed by name:

```
> data$height  
[1] 180 176 156 165
```

```
> data[, "height"]  
[1] 180 176 156 165
```

This is usually better than accessing them by column number, as the name is less likely to change than the column number.

(also, if the name changes, it will yield an error)

You can shorten the name as long as there is no ambiguity:

```
> data$height  
[1] 180 176 156 165
```

```
> data$h  
[1] 180 176 156 165
```

This is not recommended: the code may break if your script is used on a dataset that includes a new column which causes an ambiguity.

# *tidyverse is a collection of packages*



## R packages for data science

The tidyverse is an opinionated **collection of R packages** designed for data science. All packages share an underlying design philosophy, grammar, and data structures.

Install the complete tidyverse with:

```
install.packages("tidyverse")
```



## *dplyr is a grammar of data manipulation*

- Takes a data frame (or tibble) as the first argument
- Several features available:
  - mutate() adds new variables that are functions of existing variables
  - select() picks variables based on their names.
  - filter() picks cases based on their values.
  - summarise() / reframe() reduce multiple values down to a single summary.
  - arrange() changes the ordering of the rows.
- Features like connecting directly to a database, or work on data not fully loaded in memory
- Functions like bind\_rows and bind\_cols much more efficient than rbind and cbind !

## *Naive data manipulation in base R*

Take the iris data.frame, derive a new variable and calculate an average by group

```
> data(iris)
> head(iris, n = 2)
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5         1.4         0.2   setosa
2          4.9         3.0         1.4         0.2   setosa
> iris$Length.Product <- iris$Sepal.Length * iris$Petal.Length
> iris <- iris[iris$Length.Product >= 7, ]
> iris <- iris[, c("Species", "Length.Product")]
> sqldf("select Species, avg(`Length.Product`) as
avg_length_product from iris group by Species")
  Species avg_length_product
1   setosa          7.922727
2 versicolor         25.466600
3  virginica         36.873800
```

*Now with dplyr*

Take the iris data.frame, derive a new variable and calculate an average by group

```
iris.modified <-  
  iris |>  
  mutate(Length.Product = Sepal.Length * Petal.Length) |>  
  filter(Length.Product >= 7) |>  
  select(Species, Length.Product) |> # optional  
  reframe(avg.length.product = mean(Length.Product), .by =  
Species)
```

	Species	avg.length.product
1	setosa	7.922727
2	versicolor	25.466600
3	virginica	36.873800

Some useful features like everything() (see also relocate() )

```
> data(iris)
> iris <- iris |> mutate(Length.Product = Sepal.Length *
Petal.Length)
> head(iris[, 1:4], n = 2)
  Sepal.Length Sepal.Width Petal.Length Petal.Width
1           5.1          3.5          1.4          0.2
2           4.9          3.0          1.4          0.2
> iris <- iris |> select(Species, Length.Product, everything())
> head(iris[, 1:4], n = 2)
  Species Length.Product Sepal.Length Sepal.Width
1  setosa           7.14          5.1          3.5
2  setosa           6.86          4.9          3.0
```

## Or range of columns using the names

```
> head(iris |> select(Sepal.Length:Petal.Width))
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
1	5.1	3.5	1.4	0.2
2	4.9	3.0	1.4	0.2
3	4.7	3.2	1.3	0.2
4	4.6	3.1	1.5	0.2
5	5.0	3.6	1.4	0.2
6	5.4	3.9	1.7	0.4

## *Getting information about R objects:* `summary()`

The `summary()` command gives some brief information about an R object.

```
> summary(mylist)
```

	Length	Class	Mode
ages	4	-none-	numeric
height	4	-none-	numeric
sex	4	-none-	character

## *Getting information about R objects:* `summary()`

Its output depends on the type of object:

```
> summary(mylist)
```

	Length	Class	Mode
ages	4	-none-	numeric
height	4	-none-	numeric
sex	4	-none-	character

```
> summary( rnorm(100) )
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-2.62861	-0.49714	0.14360	0.05439	0.69587	2.11127

## *Getting information about R objects*

The `str()` command gives detailed information about the **structure** of an R object:

```
> str(mylist)
```

```
List of 3
```

```
$ ages : num [1:4] 21 32 41 45
```

```
$ height: num [1:4] 180 176 156 165
```

```
$ sex : chr [1:4] "M" "M" "F" "M"
```



## *Getting information about R objects*

The information provided by `str()` can indeed be really detailed; try the following commands :

```
> model <- lm( runif(10) ~ rnorm(10) )  
> str(model)
```

# Getting information about R objects

```
# Try this one if you don't believe the word "detailed" above
> model <- lm( runif(10) ~ rnorm(10) )
> str(model)
List of 12
 $ coefficients : Named num [1:2] 0.5486 0.0335
 .. attr(*, "names")= chr [1:2] "(Intercept)" "rnorm(10)"
 $ residuals    : Named num [1:10] -0.255582 -0.192832 -0.000517 0.340288 -0.336684 ...
 .. attr(*, "names")= chr [1:10] "1" "2" "3" "4" ...
 $ effects      : Named num [1:10] -1.744 -0.1102 0.0902 0.4255 -0.2811 ...
 .. attr(*, "names")= chr [1:10] "(Intercept)" "rnorm(10)" "" "" ...
 $ rank         : int 2
 $ fitted.values: Named num [1:10] 0.521 0.565 0.573 0.568 0.538 ...
 .. attr(*, "names")= chr [1:10] "1" "2" "3" "4" ...
 $ assign       : int [1:2] 0 1
 $ qr           :List of 5
 ..$ qr        : num [1:10, 1:2] -3.162 0.316 0.316 0.316 0.316 ...
 .. .. attr(*, "dimnames")=List of 2
 .. .. ..$ : chr [1:10] "1" "2" "3" "4" ...
 .. .. ..$ : chr [1:2] "(Intercept)" "rnorm(10)"
 .. .. attr(*, "assign")= int [1:2] 0 1
 ..$ qraux: num [1:2] 1.32 1.19
 ..$ pivot: int [1:2] 1 2
 ..$ tol   : num 1e-07
 ..$ rank  : int 2
 .. attr(*, "class")= chr "qr"
 $ df.residual : int 8
 $ xlevels     : Named list()
 $ call        : language lm(formula = runif(10) ~ rnorm(10))
 $ terms       :Classes 'terms', 'formula' language runif(10) ~ rnorm(10)
 .. .. attr(*, "variables")= language list(runif(10), rnorm(10))
 .. .. attr(*, "factors")= int [1:2, 1] 0 1
 .. .. .. attr(*, "dimnames")=List of 2
 .. .. .. ..$ : chr [1:2] "runif(10)" "rnorm(10)"
 .. .. .. ..$ : chr "rnorm(10)"
 .. .. attr(*, "term.labels")= chr "rnorm(10)"
 .. .. attr(*, "order")= int 1
 .. .. attr(*, "intercept")= int 1
 .. .. attr(*, "response")= int 1
 .. .. attr(*, ".Environment")=<environment: R_GlobalEnv>
 .. .. attr(*, "predvars")= language list(runif(10), rnorm(10))
 .. .. attr(*, "dataClasses")= Named chr [1:2] "numeric" "numeric"
 .. .. .. attr(*, "names")= chr [1:2] "runif(10)" "rnorm(10)"
 $ model       :'data.frame': 10 obs. of 2 variables:
 ..$ runif(10): num [1:10] 0.266 0.372 0.573 0.908 0.202 ...
 ..$ rnorm(10): num [1:10] -0.82 0.487 0.738 0.576 -0.305 ...
 .. attr(*, "terms")=Classes 'terms', 'formula' language runif(10) ~ rnorm(10)
 .. .. attr(*, "variables")= language list(runif(10), rnorm(10))
 .. .. .. attr(*, "factors")= int [1:2, 1] 0 1
 .. .. .. .. attr(*, "dimnames")=List of 2
 .. .. .. .. ..$ : chr [1:2] "runif(10)" "rnorm(10)"
 .. .. .. .. ..$ : chr "rnorm(10)"
 .. .. .. attr(*, "term.labels")= chr "rnorm(10)"
 .. .. .. attr(*, "order")= int 1
 .. .. .. attr(*, "intercept")= int 1
 .. .. .. attr(*, "response")= int 1
 .. .. .. attr(*, ".Environment")=<environment: R_GlobalEnv>
 .. .. .. attr(*, "predvars")= language list(runif(10), rnorm(10))
 .. .. .. attr(*, "dataClasses")= Named chr [1:2] "numeric" "numeric"
 .. .. .. .. attr(*, "names")= chr [1:2] "runif(10)" "rnorm(10)"
 - attr(*, "class")= chr "lm"
```

**A question...**

```
# Simulate data for 3 groups
set.seed(1)
groups <- rep( 1:3, each=10 )

measure <- vector(length=30)
measure[ groups==1 ] <- 5
measure[ groups==2 ] <- 1
measure[ groups==3 ] <- 5
measure <- measure + rnorm(30)



# Perform a one-way ANOVA on this data
boxplot( measure ~ groups )
summary( aov( measure ~ groups ) )
```

```
# Perform a one-way ANOVA on this data
```

```
> boxplot( measure ~ groups )
```

```
> summary(aov( measure ~ groups ))
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
groups	1	0.09	0.088	0.018	0.893
Residuals	28	134.85	4.816		



# Factors

Factors represent **categorical variables** in R.

They are vectors that can contain only values from a (finite) predefined set.

## *Example*

```
> hair <- factor(c("blond", "brown", "red", "blond"))
```



## *Example*

```
> hair <- factor(c("blond", "brown", "red", "blond"))
```

```
> hair
```

```
[1] blond brown red    blond
```

```
Levels: blond brown red
```

## *Example*

```
> hair <- factor(c("blond", "brown", "red", "blond"))
```

```
> hair  
[1] blond brown red    blond  
Levels: blond brown red
```

```
> hair[2] <- "blond"
```

```
> hair  
[1] blond blond red    blond  
Levels: blond brown red
```

## *Example*

```
> hair <- factor(c("blond", "brown", "red", "blond"))
```

```
> hair  
[1] blond brown red    blond  
Levels: blond brown red
```

```
> hair[2] <- "blond"
```

```
> hair[2] <- "grey"
```

```
Warning message:
```

```
In `[<-.factor`(`*tmp*`, 2, value = "grey") :  
invalid factor level, NAs generated
```

## *Example*

```
> hair <- factor(c("blond", "brown", "red", "blond"))
```

```
> hair  
[1] blond brown red    blond  
Levels: blond brown red
```

```
> hair[2] <- "blond"
```

```
> hair[2] <- "grey"
```

```
Warning message:
```

```
In `[<-.factor`(`*tmp*`, 2, value = "grey") :  
  invalid factor level, NAs generated
```

```
> hair
```

```
[1] blond <NA> red    blond  
Levels: blond brown red
```

```
> class(hair)
[1] "factor"
> typeof(hair); mode(hair)
[1] "integer"
[1] "numeric"

> as.numeric(hair)
[1] 1 NA 3 1
> as.character(hair)
[1] "blond" NA "red" "blond"
```

Internally, R stores factors as integer numbers, along with the correspondance between number and labels (1=blond, 2=brown, 3=red).

## *Ordered factors*

Use the `ordered=TRUE` option for ordinal (ordered) values:

```
> time <- factor(c(1,2,3,2,2,1), levels=c(1,2,3),  
                 labels=c("never", "sometimes", "always"),  
                 ordered=TRUE)  
  
> time  
[1] never      sometimes always      sometimes  
[5] sometimes never  
Levels: never < sometimes < always
```

## *Some R functions respect ordered factors*

Comparisons work as expected:

```
> time
[1] never      sometimes always      sometimes
[5] sometimes never
Levels: never < sometimes < always
> time[2] < time[3]
[1] TRUE
> "sometimes" < "always"
[1] FALSE
```

Some statistical modelling or plotting functions can adapt their parameters for ordered factors.

```
# Perform a one-way ANOVA on this data
```

```
> boxplot( measure ~ groups )
```

```
> summary(aov( measure ~ groups ) )
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
groups	<b>1</b>	0.09	0.088	0.018	<b>0.893</b>
Residuals	28	134.85	4.816		

```
> groups <- as.factor(groups)
```

```
> groups
```

```
[1] 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2 3 3 3 3 3 3 3 3 3
```

```
Levels: 1 2 3
```

```
> summary(aov( measure ~ groups ) )
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
groups	<b>2</b>	94.12	47.06	52.95	<b>4.53e-10 ***</b>
Residuals	27	24.00	0.89		

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```



Recent versions of R have introduced several major changes with regards to how factors are handled.

## *Concatenating factors: before R 4.1*

```
> c(hair, hair)
[1] 1 2 3 1 1 2 3 1

# Workaround #1
> factor( as.character(hair), as.character(hair) )

# Workaround #2
> unlist( list( hair, hair) )
```

Simply concatenating factors was creating a vector made out of the numeric values, which is almost certainly not what you want.

This behaviour has changed starting in R version 4.1  
(May 2021)

### Changelog:

Using `c()` to combine a factor with other factors now gives a factor, and specifically an ordered factor when combining ordered factors with identical levels.

In R versions before 4.0.0, by default, `data.frame()` and `read.table()` convert all non-numerical values into factors.

This can be useful, or (more often...) it can be annoying.

Options to change this behaviour:

- `stringsAsFactors=FALSE`, or
- `as.is=TRUE` (for `read.table` only)

It can also be set by default using

```
options(stringsAsFactors=FALSE)
```

But this is not recommended, as your code may not work anymore if someone else uses it without specifying the same default option.

## CHANGES IN R 4.0.0

### SIGNIFICANT USER-VISIBLE CHANGES

R now uses a `stringsAsFactors = FALSE` default, and hence by default no longer converts strings to factors in calls to `data.frame()` and `read.table()`.

A large number of packages relied on the previous behaviour and so have needed/will need updating.

## *Factors and memory size*

In old versions of R, using factors for long vectors could save memory :

```
> f1 <- sample( c("Homo Sapiens", "Mus Musculus"), 10000,
                replace=TRUE)
> summary(f1)
  Length      Class      Mode 
 10000 character character
> table(f1)
f1
Homo Sapiens      Mus Musculus 
      4945             5055 
> object.size(f1)
80168 bytes
> f2 <- factor(f1)
> object.size(f2)
40544 bytes
```

## *Factors and memory size*

In recent versions of R (2.6+) it is not the case anymore, as R stores only once each occurrence of a string in a vector:

```
> f1 <- sample( c("Homo Sapiens", "Mus Musculus"), 10000,
                replace=TRUE)
> summary(f1)
      Length      Class      Mode 
    10000 character character 
> table(f1)
f1
Homo Sapiens      Mus Musculus 
      4945             5055 
> object.size(f1)
40104 bytes
> f2 <- factor(f1)
> object.size(f2)
40312 bytes
```

**What do these commands do ?**

**c=c (c=c)**

**c=c (c="c")**



**What do these commands do ?**

**c=c (c=c)**

**Let's read from right to left**

**c=c (c=c)**

**Let's read from right to left**

**c=c (c=c)**

c=c (c=c)

```
> c
```

```
function (...) .Primitive("c")
```

# The `c()` function, as an object

`c=c (c=c)`

```
> c  
function (...) .Primitive("c")
```

Let's read (more or less)  
from right to left

c=c (c=c)

Let's read (more or less)  
from right to left

c=c (c=c)

Calling the `c()` function,  
in order to create a vector

`c=c(c=c(c))`



**Calling the `c()` function,  
in order to create a vector**

`c=c(c=c)`



The vector contains a single element:  
the “c” object

**c=c (c=c)**

c=c (c=c)

**Giving the name «c» to  
the element of the vector**

**c=c (c=c)**

`c=c (c=c)`



# Storing the vector in a new «c» object

The diagram illustrates the process of storing a vector in a new object. It consists of two parts: a variable assignment and a function call. In the first part, a red circle containing the letter 'c' is followed by an equals sign and a green circle containing the letter 'c'. In the second part, there is an opening parenthesis, followed by a green circle containing 'c', an equals sign, another green circle containing 'c', and a closing parenthesis. This represents the code 'c = c(c)'.

# The final result

**c=c (c=c)**

```
> c=c (c=c)
```

```
> c
```

```
$c
```

```
function (...) .Primitive("c")
```



# What about the other one ?

**c=c (c="c")**

```
> c=c (c="c")
```

```
> c
```

```
c
```

```
"c"
```

**One more question...**

# One more question...

**How does R store both the `c()` function and the `c` vector, and how does it differentiate between them ?**

```
> c=c (c=c)
> c=c (c="c")
> c
  c
"c"
```

**We will come back to this later !**

## *Reminder: getting information about R objects*

The `summary()` command gives some brief information about an R object; its output depends on the type of object:

```
> summary(mylist)
```

	Length	Class	Mode
ages	4	-none-	numeric
height	4	-none-	numeric
sex	4	-none-	character

```
> summary(aov( measure ~ groups ) )
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
groups	<b>1</b>	0.09	0.088	0.018	<b>0.893</b>
Residuals	28	134.85	4.816		

*Reminder: getting information about R objects*

## How does `summary()` know what to print for different objects ?

```
> summary(mylist)
```

	Length	Class	Mode
ages	4	-none-	numeric
height	4	-none-	numeric
sex	4	-none-	character

```
> summary(aov( measure ~ groups ) )
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
groups	<b>1</b>	0.09	0.088	0.018	<b>0.893</b>
Residuals	28	134.85	4.816		